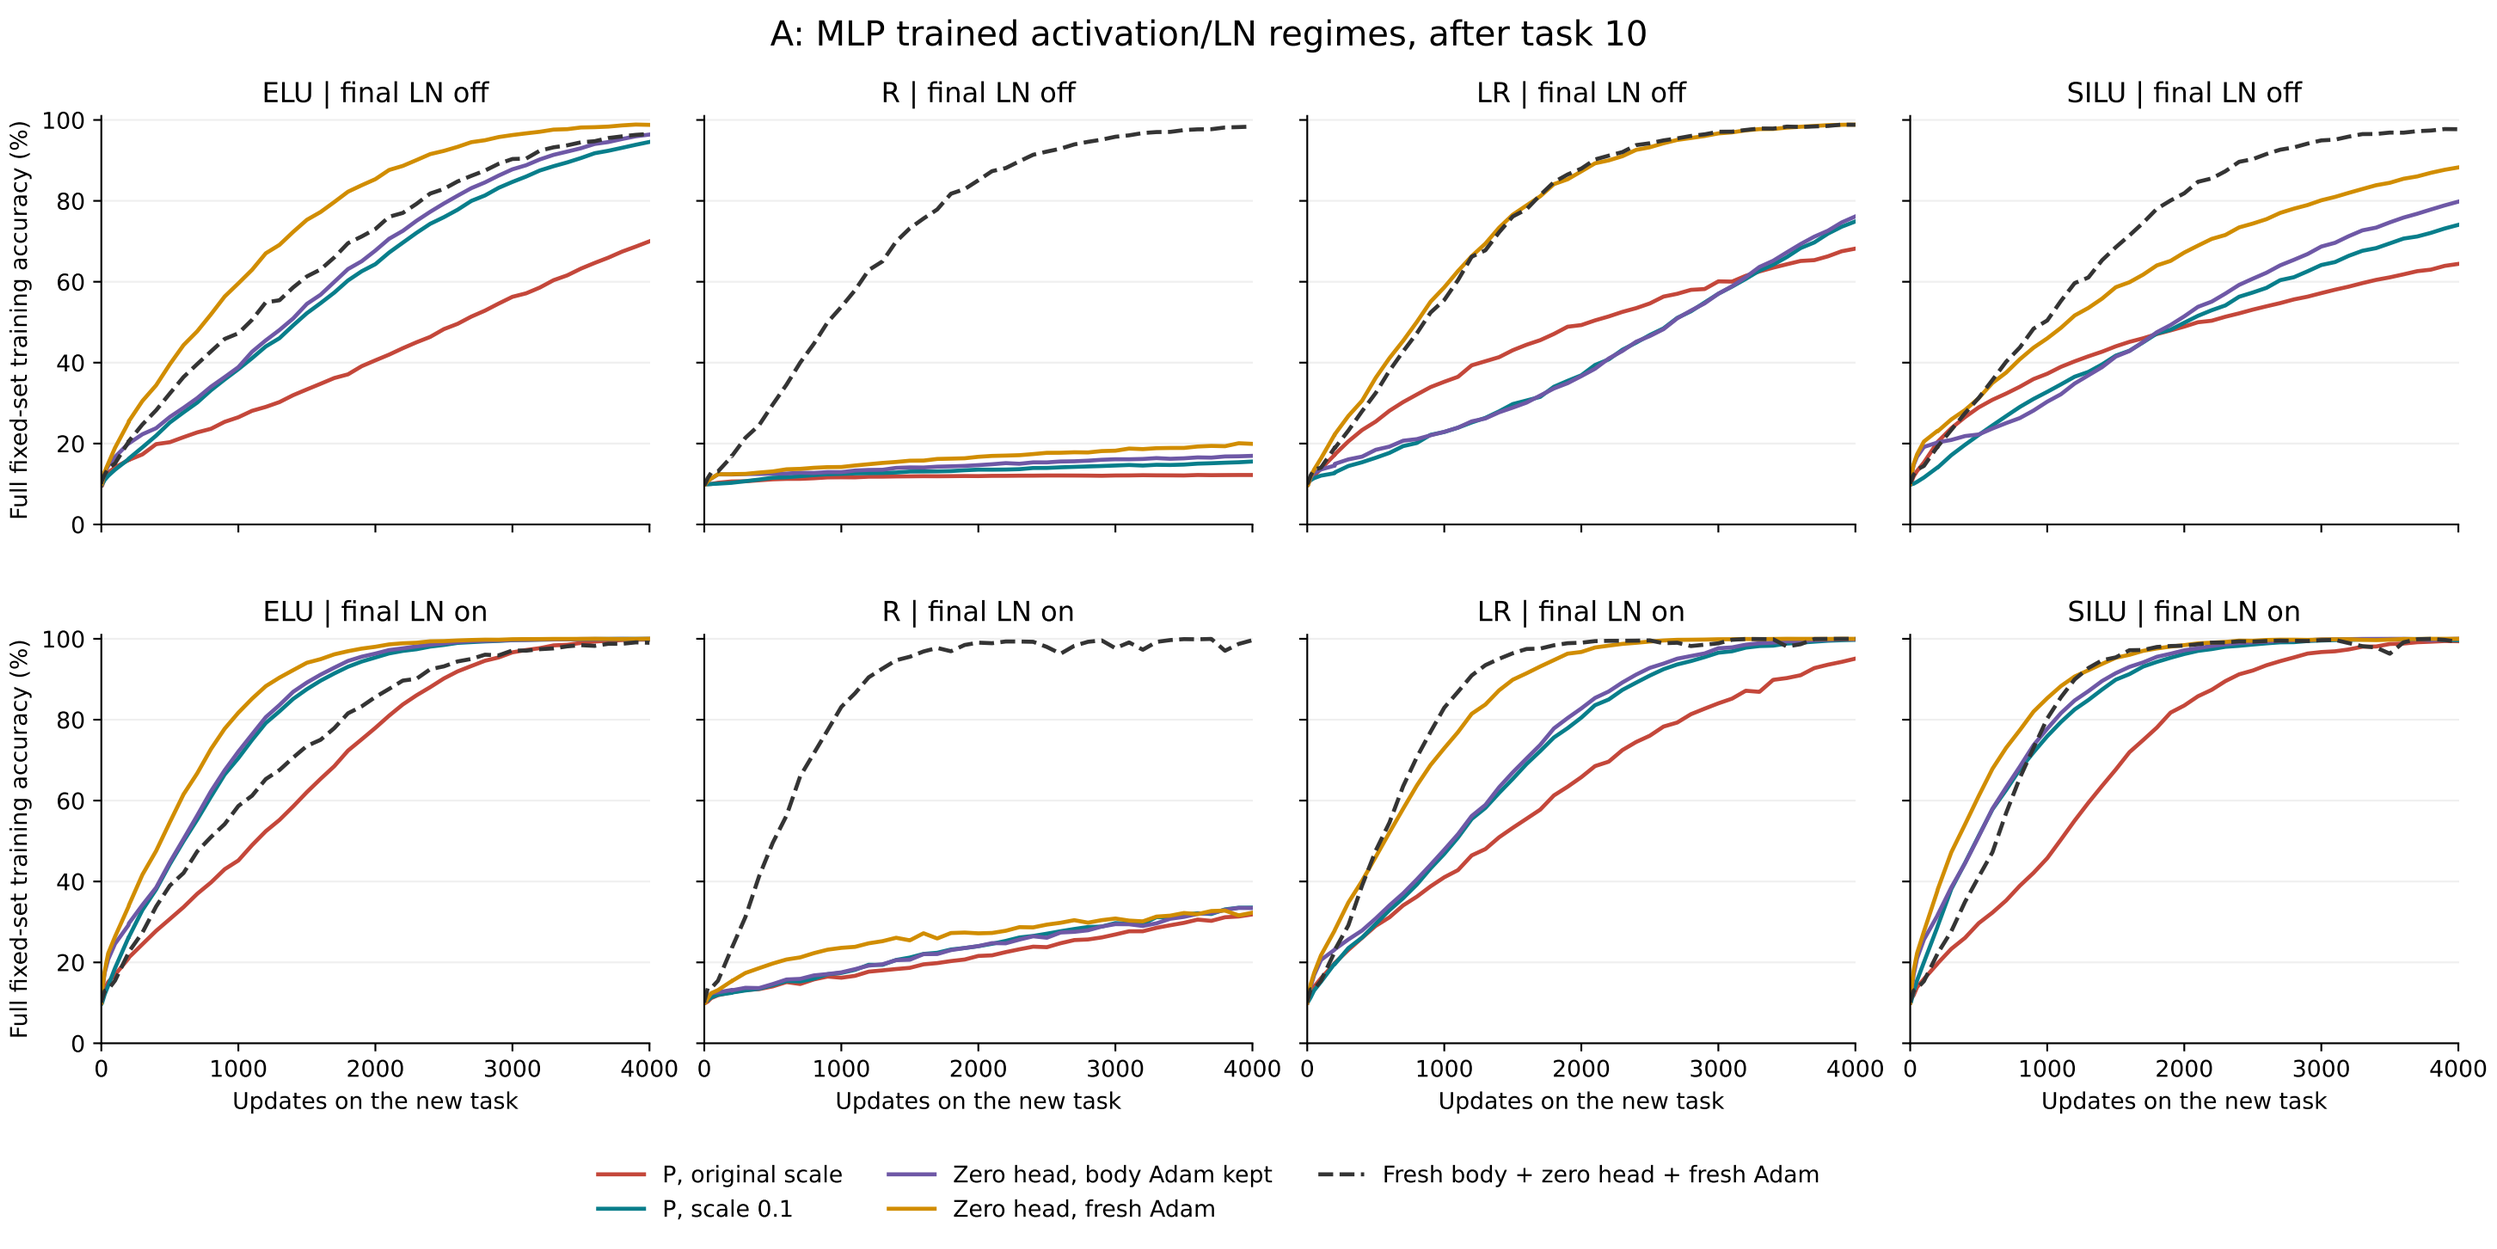
A: MLP trained activation/LN regimes, after task 10
ELU | final LN off
R | final LN off
LR | final LN off
SILU | final LN off
ELU | final LN on
R | final LN on
LR | final LN on
SILU | final LN on
Full fixed-set training accuracy (%)
0
20
40
60
80
100
0
1000
2000
3000
4000
Updates on the new task
P, original scale
P, scale 0.1
Zero head, body Adam kept
Zero head, fresh Adam
Fresh body + zero head + fresh Adam
Mean of 10 paired seeds. One-time head edits; all parameters learn. Fresh body uses paired initialization, labels and batches.